

---

# SkiLD: Unsupervised Skill Discovery Guided by Local Dependencies

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

1        Unsupervised skill discovery carries the promise that an intelligent agent can learn  
2        reusable skills through autonomous, reward-free interactions with environments.  
3        Existing unsupervised skill discovery methods learn skills by encouraging distin-  
4        guishable behaviors that cover diverse states. However, in complex environments  
5        with many state factors (e.g., household environments with many objects), learning  
6        skills that cover all possible states is impossible, and naively encouraging state  
7        diversity often leads to simple skills that are not ideal for solving downstream  
8        tasks. This work introduces Skill Discovery from Local Dependencies (SkiLD),  
9        which leverages state factorization as a natural inductive bias to guide the skill  
10       learning process. The key intuition guiding SkiLD is that skills that induce **diverse**  
11       **interactions** between state factors are often more valuable for solving downstream  
12       tasks. To this end, SkiLD develops a novel skill learning objective that explicitly  
13       encourages the mastering of skills that effectively induce different interactions  
14       within an environment. We evaluate SkiLD in several domains with challenging,  
15       long-horizon sparse reward tasks including a realistic simulated household robot  
16       domain, where SkiLD successfully learns skills with clear semantic meaning and  
17       shows superior performance compared to existing unsupervised reinforcement  
18       learning methods that only maximize state coverage.

## 19    1 Introduction

20    Reinforcement learning (RL) achieves impressive successes when solving decision-making problems  
21    with well-defined reward functions [62, 19, 31]. However, creating this reward function is often not  
22    trivial [6]. In contrast, humans and other intelligent creatures can learn, without external reward  
23    supervision, behaviors that produce repeatable and predictable changes in the environment [17].  
24    These behaviors, which we call *skills*, can be later repurposed to solve downstream tasks efficiently.  
25    One of the promises of such “unsupervised RL” is to endow artificial agents with similar capabilities  
26    to discover reusable skills without explicit rewards.

27    One predominant strategy of prior skill discovery methods focuses on training skills to reach diverse  
28    states while being distinguishable [18, 57, 48]. However, in complex environments that contain many  
29    *state factors*—distinct elements such as individual objects in a household (a formal description in  
30    Sec. 2.1), the exponential number of distinct states makes it impossible to learn skills that cover every  
31    state. Consequently, these methods result in simple skills that only change the easy-to-control factors  
32    (e.g., in a manipulation task moving the agent itself to diverse positions or manipulating each factor  
33    independently), and fail to cover other desirable but challenging behaviors. Unsurprisingly, these  
34    simple skills often struggle to solve meaningful tasks, resulting in poor downstream performance.

35    Our key insight is that, given a factored state space, the interactions between state factors can often  
36    act as a powerful inductive bias in guiding the learning of useful skills. For example, in a household




Figure 1: **Skill Discovery from Local Dependencies (SkiLD)** describes skills that encode interactions (i.e., local dependencies) between state factors. In contrast to prior diversity-based methods that can easily get stuck by moving the robot to diverse, but non-interactive states, and factor-based methods that are trained to manipulate the hammer and nail, but not their interactions, SkiLD not only manipulate each object (left, middle) but also induce interactions between them (right), by specifying different local dependencies. These skills are often more useful than the “easy” skill learned by previous methods for downstream task-solving.

environment, a skill that induces interactions between a knife and a fruit is more likely to encode fruit-cutting behaviors that can be crucial for a wide range of downstream kitchen tasks. Furthermore, exploring state coverage within the space of interaction states can uncover desirable interactions like cutting the peach.

To this end, we introduce Skill Discovery from Local Dependencies (SkiLD), a novel skill discovery method that explicitly learns skills that induce diverse interactions. Specifically, SkiLD models the interactions between state factors using the framework of *local dependencies* (where local refers to state-specific, see details in Sec. 2.2) and proposes a novel intrinsic reward that 1) encourages the agent to induce specified interactions, and 2) encourages the agent to discover diverse ways of inducing specified interaction, as visualized in Figure 2. During skill learning, SkiLD gradually discovers new interactions and learns to induce them, based on the skills that it already mastered, resulting in a diverse set of interaction-inducing behaviors that can be readily repurposed for downstream tasks. During task learning, the skill policy is reused, and a task-specific policy is learned to select (a sequence of) skills to maximize task rewards efficiently.

We evaluate the performance of SkiLD on factor-rich environments with 10 downstream tasks against existing unsupervised reinforcement learning methods. Our experiments indicate that SkiLD learns to induce diverse interactions and outperforms other methods on most of the examined tasks.

## 2 Background

In this paper, our unsupervised skill discovery method is set up in a factored Markov decision process and builds off previous diversity-based methods, as described in Sec. 2.1. To enhance the expressivity of skills, our method further augments the skill representation with interactions between state factors, which we formalize as local dependencies as described in Sec. 2.2.

### 2.1 Factored Markov Decision Process (Factored MDP)

We consider unsupervised skill discovery in a reward-free Factored Markov Decision Process [7] defined by the tuple  $\mathcal{M} = (\mathcal{S}, \mathcal{A}, p)$ .  $\mathcal{S} = \mathcal{S}^1 \times \dots \times \mathcal{S}^N$  is a factored state space with  $N$  subspaces, where each subspace  $\mathcal{S}^i$  is a multi-dimensional continuous or discrete random variable. Then, correspondingly, each state  $s \in \mathcal{S}$  consists of  $N$  state factors, i.e.,  $s = (s^1, \dots, s^N)$ ,  $s^i \in \mathcal{S}^i$ . In this paper, we use uppercase letters to denote random variables and lowercase for their specific values (e.g.,  $S$  denotes the random variable for states  $s$ ).  $\mathcal{A}$  is the action space, and  $p$  is an unknown Markovian transition model that captures the probability distribution over the next state  $S' \sim p(\cdot | S, A)$ .

The factorization in  $\mathcal{S}$  inherently exists in many environments, and is a common assumption in prior unsupervised skill discovery works [21, 27]. For example, in robotics, an environment typically consists of a robot and several objects to manipulate, and, for each object,  $\mathcal{S}^i$  would represent its attributes of interest, like pose. In this work, we explore, when the state factorization is given, how we can utilize it to improve unsupervised skill discovery. In practice, the factorization can either be directly provided by the environment or obtained from image observations with existing disentangled representation learning methods [45, 29].

Following prior work, our method consists of two stages—skill learning and task learning. During the skill learning phase, we seek to learn a skill policy  $\pi_\omega(\cdot|s, z)$ , which defines a conditional distribution over actions given the current state  $s$  and some skill representation  $z$ , where skills indicate the desired behaviors of the agent. Once the skills are learned, they can be chained together to solve downstream tasks during the task learning phase through an extrinsic reward-optimizing policy. During task learning, a downstream task reward function  $r : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$  is provided by the environment. A high-level policy  $\pi(z|s)$  is then trained to optimize the expected return through outputting correct skills  $z$  given state  $s$ .

## 2.2 Identifying Local Dependencies between State Factors

A key insight of SkiLD is to utilize interactions (or, formally, local dependencies) between state factors as part of the skill representation. In this section, we first formally define local dependencies, introduce their identification, and finally discuss their application to factored MDPs.

SkiLD takes a causality-inspired approach for defining and detecting local dependencies [5, 56], where we use *local* to refer to a particular assignment of values for a random variable, as opposed to *global* which applies to all values. Formally, for an *event of interest*  $Y$  and its potential causes  $X = (X^1, \dots, X^N)$ , given the value of  $X = x$ , local dependencies focus on which  $X^i$ 's are the state-specific cause of the outcome event  $Y = y$  (for simplicity of presentation, in this section we overload  $N$  as the number of potential causes rather than number of variables and  $p$  as the transition function according to a subset of the variables). Formally, we denote the general data generation process of  $Y$  as  $p : X \rightarrow Y$  and the data generation process when  $Y$  is *only influenced* by a subset of  $X$  as  $p^{\bar{X}} : \bar{X} \rightarrow Y$ , where  $\bar{X} \subseteq X$ . Then, given the value of all variables,  $X^1 = x^1, \dots, X^N = x^N$  and  $Y = y$ , we say  $Y$  locally depends on  $\bar{X}$ , if  $\bar{X}$  is the *minimal* subset of  $X$  such that knowing their values is necessary and sufficient to generate the result of  $Y = y$ , i.e.,

$$\arg \min_{\bar{X} \subseteq X} |\bar{X}| \quad \text{s.t.} \quad p^{\bar{X}}(Y = y | \bar{X} = \bar{x}) = p(Y = y | X = x), \quad (1)$$

where  $|\bar{X}|$  is the number of variables in  $\bar{X}$ . For example, suppose that a robot opens a refrigerator door in a particular transition. The event of interest  $Y$  is the refrigerator door becoming open, and it locally depends on two factors: the robot and the refrigerator door, while other state factors such as objects inside the refrigerator do not locally influence  $Y$ .

To identify local dependencies, one can conduct a conditional independence test  $y \perp\!\!\!\perp x^i | \{x/x^i\}$  to examine whether a variable  $X^i$  is necessary for predicting  $Y = y$ . In prior works, one form of this test is to examine whether the pointwise conditional mutual information (pCMI) is greater than 0,

$$\text{pCMI}(y; x^i | \{x/x^i\}) = \log \frac{p(y|x)}{p^{X/X^i}(y | \{x/x^i\})} > 0. \quad (2)$$

If so, then it suggests that knowing  $X^i = x$  provides additional information about  $Y$  that is not present in  $\{X/X^i\}$ , and  $Y$  locally depends on  $X^i$ . As the data generation processes are generally unknown, one has to approximate them with learned models. Recent work in RL has utilized various approximations such as attention weights [51], Granger causality [13], and input gradients [60].

In this work, for a transition ( $\mathcal{S} = s, \mathcal{A} = a, \mathcal{S}' = s'$ ), the event of interest is each next state factor being  $(\mathcal{S}^i)' = (s^i)'$ , and we infer whether it locally depends on each state factor  $\mathcal{S}^j$  and the action  $\mathcal{A}$  (i.e., whether there is an interaction between state factors  $i$  and  $j$ , where factor  $j$  influences  $i$ ). Then we aggregate all local dependencies into a state-specific dependency graph (abbreviated in this work to *dependency graph*) to enhance skill representation, as explained in detail in Section 3.

## 3 Skill Discovery from Local Dependencies (SkiLD)

In this section, we describe SkiLD, which enhances the expressivity of skills using local dependencies. SkiLD represents local dependencies as *state-specific dependency graphs*, defined in Sec. 2.2. Unlike previous unsupervised skill discovery methods that randomly sample the skill vector  $z$  from fixed distributions during skill learning, SkiLD requires a procedure to intelligently generate target dependency graphs during training. As such, SkiLD frames unsupervised skill discovery as a hierarchical RL problem, where a graph-conditioned skill policy learns to induce different local dependencies

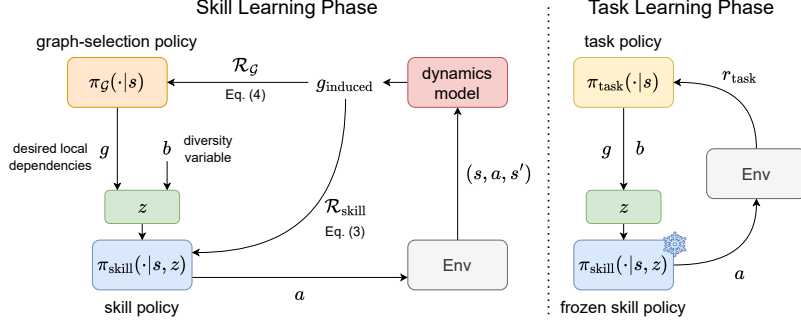


Figure 2: During **skill learning** of SkiLD, the graph-selection policy specifies desired local dependencies for the skill policy to induce, and the induced dependency graph is identified by the dynamics model and used to update both policies. During **task learning** (right), the skill policy is kept frozen and a task policy is trained to select skills to maximize task reward.

120 using primitive actions, and a high-level graph selection policy chooses which local dependencies the  
 121 skill policy should induce next to guide exploration and skill-policy learning.

122 This requires formalizing two components: (1) the skill representation  $\mathcal{Z}$  for the skill policy  
 123  $\pi_{\text{skill}}(a|s, z)$  and its corresponding reward function  $\mathcal{R}_{\text{skill}}$ , presented in Sec. 3.1, and (2) the graph  
 124 selection policy  $\pi_{\mathcal{G}}(g|s)$  and its reward function  $\mathcal{R}_{\mathcal{G}}$ , presented in Sec. 3.2.

### 125 3.1 Skill Policy

126 Prior unsupervised skill discovery methods usually focus skill learning on changing the state or each  
 127 factor diversely. Consequently, they are typically limited to learning simple skills, for example, only  
 128 changing the easiest-to-control factor in the state (i.e., the agent itself). To address this problem,  
 129 SkiLD not only focuses on changing the state but also considers the interactions between state factors.

130 **Skill Representation.** SkiLD represents the skill space as the combination of two components:  
 131  $\mathcal{Z} = \mathcal{G} \times \mathcal{B}$ , where  $g \in \mathcal{G}$  is a state-specific dependency graph that specifies the *desired* local  
 132 dependencies between state factors, and  $b \in \mathcal{B}$  is a diversity variable the same as that used in  
 133 Eysenbach et al. [18]. Together  $z \in \mathcal{Z}$  guides the agent to change the state distinguishably while  
 134 inducing particular local dependencies. Specifically, the dependency graph is represented as a binary  
 135 matrix  $\mathcal{G} = \{0, 1\}^{N \times (N+1)}$ , where each edge  $\mathcal{G}^{ij}$  denotes, during the transition  $(s, a, s')$ , whether the  
 136 state factor  $(s^i)'$  locally depends on  $s^j$ . The diversity variable  $\mathcal{B}$  can be either discrete or continuous.  
 137 In this work, without loss of generality, we use a discrete space of  $\{1, \dots, K\}$  where  $K$  is a predefined  
 138 number. During skill training, we sample the diversity variable  $b$  from a fixed uniform distribution  
 139  $p(b)$ , following the procedure of Eysenbach et al. [18].

140 Given this skill space, SkiLD learns skills as a skill-conditioned policy  $\pi_{\text{skill}} : \mathcal{S} \times \mathcal{Z} \rightarrow \mathcal{A}$ , where  
 141 the  $\pi_{\text{skill}}$  is trained to reach diverse states while ensuring that the local dependencies specified by  
 142 the graph are induced. During skill learning, we select actions by iteratively calling the skill policy  
 143  $\pi_{\text{skill}}$ , and we denote  $g_{\text{induced}}$  as the graph that describes the local dependencies induced in a transition  
 144  $(s, a, s')$  when executing a selected action  $a$ . We design the reward function of SkiLD as:

$$\mathcal{R}_{\text{skill}} = \mathbb{1}[g_{\text{induced}} = g] \cdot (1 + \lambda \mathcal{R}_{\text{diversity}}), \quad (3)$$

145 where  $\mathbb{1}[g_{\text{induced}} = g]$  measures whether the induced dependency graph matches the desired graph,  
 146  $\mathcal{R}_{\text{diversity}}$  is the weighted diversity reward that further encourages visiting diverse states when the  
 147 desired graph is induced, and  $\lambda$  is the coefficient of diversity reward. In the following paragraphs, we  
 148 describe how we infer  $g_{\text{induced}}$  and estimate  $\mathcal{R}_{\text{diversity}}$  for each transition.

149 **Inferring Induced Graphs.** To infer the induced graph for a transition  $(\mathcal{S} = s, \mathcal{A} = a, \mathcal{S}' = s')$ ,  
 150 we need to determine, for each  $(s^i)'$ , whether it locally depends on each factor  $s^j$  and the action  
 151  $\mathcal{A}$ . Specifically, following Sec. 2.2, we evaluate the conditional dependency  $(s^i)' \not\perp\!\!\!\perp s^j | \{s/s^j, a\}$   
 152 by examining whether their pointwise conditional mutual information (pCMI) is greater than a  
 153 predefined threshold  $\text{pCMI}^{ij} = \frac{p((s^i)' | \{s, a\})}{p((s^i)' | \{s/s^j, a\})} \geq \epsilon$ . If so, it suggests that  $s^j$  is necessary to predict  
 154  $(s^i)'$  and thus the local dependency exists. Meanwhile, as the transition probability  $p$  is unknown, we  
 155 approximate it with a learned dynamics model that is trained to minimize prediction error.

Finally, after obtaining the induced dependency graph, we evaluate  $\mathbb{1}[g_{\text{induced}} = g]$  by examining whether each edge  $g_{\text{induced}}^{ij}$  matches the corresponding edge in the desired graph  $g^{ij}$ . As  $\mathcal{R}_{\text{skill}}$  only provides sparse rewards to the skill policy when the desired graph is induced, we use hindsight experience replay [1] to enrich learning signals, by relabelling induced graphs as desired graphs in some episodes.

**Diversity Rewards.** When the skill policy induces the desired graph,  $\mathcal{R}_{\text{diversity}}$  further encourages it to visit different distinguishable states under different diversity indicators  $b$ , e.g., driving the nail to different locations. This diversity enhances the applicability of learned skills. To this end, we design the diversity reward  $\mathcal{R}_{\text{diversity}}$  as the forward mutual information between visited states and the diversity indicator  $I(s; b)$ , following DIAYN. To estimate the mutual information, we approximate it with a variational lower bound  $I(s; b) \geq q(b|s)$ , where  $q(b|s)$  is a neural network discriminator trained to predict the diversity indicator  $b$  from the visited state. In practice, rather than learning a single low-level skill to handle all graphs and diversity parameters, we utilize a factorized lower-level policy, where there is a separate policy for each factor. More details about this subdivision can be found in Appendix A.

### 3.2 Graph-Selection Policy

To acquire skills that are useful for downstream tasks,  $\pi_{\text{skill}}$  needs to learn to induce a wide range of local dependencies *sample-efficiently*. To this end, we propose to learn a graph-selection policy  $\pi_{\mathcal{G}} : \mathcal{S} \rightarrow \mathcal{G}$  to guide the training of  $\pi_{\text{skill}}$ . Specifically, training  $\pi_{\text{skill}}$  requires a wise selection of graphs — as graph space  $\mathcal{G}$  increases super-exponentially in the number of state factors  $N$ , many graphs are not inducible. To this end, we only select target graphs for skill policy from a history of all seen graphs. As the agent learns to induce existing graphs in diverse ways, new graphs may be encountered, gradually expanding the set of seen graphs.

However, though this history guarantees graph inducibility, two challenges still remain: (1) How to efficiently explore novel local dependencies, especially hard-to-visit ones? (2) For all seen graphs, which one should  $\pi_{\text{skill}}$  learn next to maximize training efficiency? We address these challenges based on the following heuristic — compared to well-learned skills,  $\pi_{\text{skill}}$  should focus its training on underdeveloped skills. Meanwhile, learning new skills opens up the possibility of visiting novel local dependencies, e.g., learning to grasp the hammer makes it possible for the robot to hammer the nail.

According to this heuristic, we learn a graph-selection policy  $\pi_{\mathcal{G}}$  that guides the exploration and training of the skill policy  $\pi_{\text{skill}}$ . Specifically,  $\pi_{\mathcal{G}} : \mathcal{S} \rightarrow \mathcal{G}$  selects a new dependency graph the skill policy should induce for the next  $L$  time steps. To increase the likelihood of visiting hard graphs,  $\pi_{\mathcal{G}}$  is trained to maximize the following graph novelty reward

$$\mathcal{R}_{\mathcal{G}} = \frac{1}{\sqrt{C(g_{\text{visited}})}}, \quad (4)$$

where  $C(g_{\text{visited}})$  is the number of times that we have seen the graph in the collected transition. While Eq. 4 is similar to state-count-based exploration reward, here, it is based on the count of dependency graphs, and thus applicable to both discrete and continuous state space.

### 3.3 Downstream Task Learning

In SkiLD, we utilize hierarchical RL to solve reward-supervised downstream tasks with the discovered skills. The skill policy,  $\pi_{\text{skill}}$  acts as the low-level policy while a task policy,  $\pi_{\text{task}} : \mathcal{S} \rightarrow \mathcal{Z}$ , is learned to select which skill  $z = (g, b)$  to execute for  $L$  steps. Compared to diversity-based skills that are limited to simple behaviors, our local-dependency-based skills enable a wide range of interactions between state factors, leading to more efficient exploration and superior performance of downstream tasks learning.

## 4 Experiments

In this section we aim to provide empirical evidence towards the following questions: **Q1)** Do the skills learned by SkiLD induce a diverse set of interactions among state factors? **Q2)** Do the skills learned by SkiLD enable more efficient downstream task learning compared to other




Figure 3: **Evaluation domains:** Mini-behavior: Installing Printer, Thawing and Cleaning Car, and iGibson.

203 unsupervised reinforcement learning methods? Our learned skills can be visualized at <https://sites.google.com/view/skild/>.  
 204

#### 205 4.1 Domains

206 In this work, we focus on addressing the challenge of vast state space brought by a large number  
 207 of state factors. Hence, we evaluate our method on two challenging *object-rich* embodied AI  
 208 benchmarks: Mini-behavior [30] and Interactive Gibson [40].

209 The **Mini-behavior (Mini-BH) domain** [30] (Figure 3a) contains a set of gridworld environments  
 210 where an agent can move around and interact with a variety of objects to accomplish certain household  
 211 tasks. While conceptually simple, this domain has been shown to be extremely challenging for Vanilla  
 212 RL with sparse reward [30]. Each Mini-BH environment contains different objects and different  
 213 success criteria. We tested on three particular environments in Mini-behavior, including:

- 214 • **Installing Printer:** A relatively simple environment with three state factors: the agent, a table, and  
 215 a printer that can be installed.
- 216 • **Cleaning Car:** An environment where the objects have rich and complex interactions. The state  
 217 factors include the agent, a toggleable sink, a piece of rag that can be soaked in the sink, a car that  
 218 the rag can clean, a soap and a bucket which can together be used to clean the rag.
- 219 • **Thawing:** An environment with lots of movable objects. The state factors include the agent, a sink,  
 220 a fridge that can be opened, and three objects that can be thawed in the sink: fish, olive, and a date.

221 The **Interactive Gibson (iGibson) domain** [41] (Figure 3b) contains a realistic simulated Fetch  
 222 Robot that operates in a kitchen environment with a refrigerator, sink, knife, and peach. The peach  
 223 can be washed or cut. This domain is very difficult especially when using low-level motor commands  
 224 because much of the domain is free space, meaning that only a minute fraction of action sequences  
 225 will manipulate the objects meaningfully.

226 Both Mini-BH and iGibson require learning long-horizon policies spanning many low-level actions  
 227 from sparse reward, making these challenging domains (see details in Appendix).

#### 228 4.2 Baselines

229 Before evaluating the empirical questions, we provide a brief description of the baselines. These  
 230 baselines include unsupervised skill learning, and causal and hierarchical methods.

231 **Diversity is all you need (DIAYN [18]):** This method learns unsupervised state-covering skills using  
 232 a mutual information objective. SkiLD utilizes a version of this for state-diversity skills modulated by  
 233 a desired dependency graph. This baseline determines how incorporating graph information affects  
 234 the algorithm.

235 **Controllability-Aware Skill Discovery (CSD [48]):** Extends DIAYN with a factorization based on  
 236 controllability. This baseline is a comparable skill learning method that leverages state factorization  
 237 but does not encode local dependencies.

238 **Exploration via Local Dependencies (ELDEN [60]):** This method utilizes gradient-based techniques  
 239 to infer local dependencies for exploration. However, without a skill learning component, it can  
 240 struggle to chain together complex behavior.

**Chain of Interaction Skills (COInS [13]):** This is a hierarchical algorithm that constructs a chain of skills using Granger-causality to identify local dependencies. Because it is restricted to pairwise interactions, it struggles to represent the rich policies necessary for these tasks.

**Vanilla RL:** This baseline uses PPO [55] to directly train an agent with the extrinsic reward. Unlike other baselines, this method does not have a pertaining phase. Since all the task rewards are sparse and the tasks are often long horizon, vanilla RL often struggles.

### 4.3 Interaction Graph Diversity

We first evaluate whether SkiLD is indeed capable of achieving complex interaction graphs (Q1), comparing against two strong skill discovery baselines introduced earlier: DIAYN and CSD.

Each of these methods is trained for 10 Million steps without having access to any reward. Then to evaluate their learned skills, we unroll each of them for 500 episodes with randomly sampled skills  $z$  and examine the diversity of the interaction graphs they can induce. Figure 4 illustrates the percentages of episodes where particular local dependencies have been induced at least once, in Mini-BH Cleaning Car. We find that DIAYN and CSD are limited to skills that only manipulate one object individually, i.e. (agent, rag, action  $\rightarrow$  rag) or (agent, soap, action  $\rightarrow$  soap). By contrast, SkiLD learns to induce more complicated causal interactions, such as soaking the rag in the sink (sink, rag  $\rightarrow$  rag) and cleaning the car with the soaked mug (car, rag  $\rightarrow$  car).

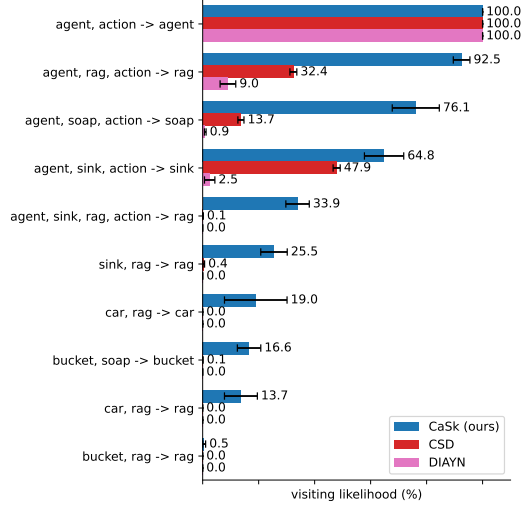


Figure 4: The percentage of episodes where a dependency graph is induced through random skill sampling. Standard deviation is calculated across five random seeds.

### 4.4 Performance

Next, we evaluate whether the local dependency coverage provided by SkiLD leads to a performance boost in downstream task learning (Q2). We follow the evaluation setup in the unsupervised reinforcement learning benchmark [36], where for a given environment, an agent is first pre-trained without access to task reward for  $K_{pt}$  steps, and then finetuned for  $K_{ft}$  steps. Importantly, the same pre-trained skills are reused on multiple distinct downstream tasks within the same environment, so that only the upper-level skill-selection policy is task-specific. We have  $K_{pt} = 2M$ ,  $K_{ft} = 1M$  for installing printer,  $K_{pt} = 10M$ ,  $K_{ft} = 5M$  for thawing and cleaning car, and  $K_{pt} = 4M$ ,  $K_{ft} = 2M$  for iGibson, and evaluate each method for each task across 5 random seeds. Hyperparameter details can be found in Appendix D. Specifically, we evaluate on the following downstream tasks:

- **Installing Printer:** We have a single downstream task in this environment, where the agent needs to pick up the printer, put it on the table, and turn it on.
- **Thawing:** We have three downstream tasks: thawing the fish or the olive or the date.
- **Cleaning Car:** We consider three downstream tasks, where each task is a pre-requisite of the following one. The tasks are: soak the rag in the sink; clean the car with the rag; and clean the dirty rag using the soap in the bucket.
- **iGibson:** The tasks for this domain are: grasping the peach, washing the peach in the sink, and cutting the peach with a knife.

After skill learning, we train a new upper-level policy that uses  $z$  as actions and is trained with extrinsic reward, as described in Section 3.3. Figure 5 illustrates the improvement of SkiLD as compared to other methods. Without combining dependency graphs with skill learning, other methods struggle with any but the simpler tasks. COInS performs poorly because of its chain structure, which restricts the agent controlling policy from picking up objects. ELDEN’s exploration reaches graphs, but

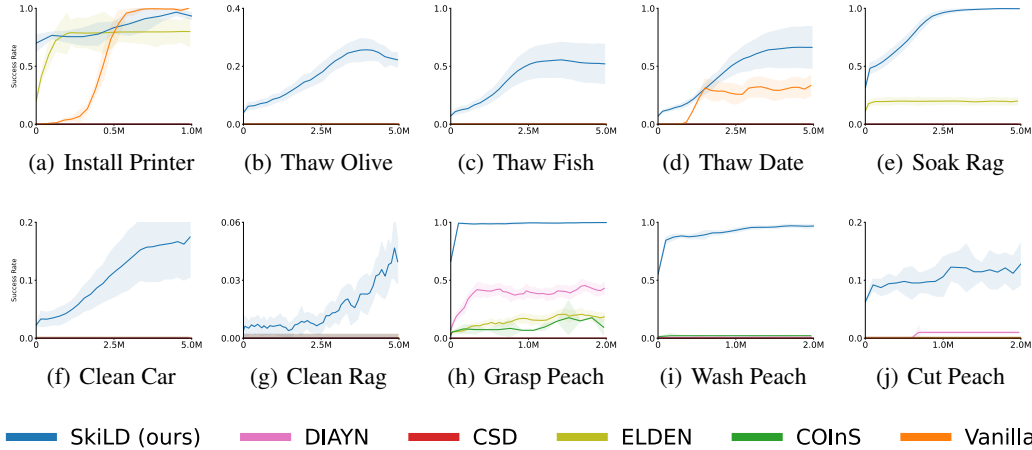


Figure 5: Training curves of SkiLD and baselines on multiple downstream tasks (reward supervised second phase). Each curve depicts the mean and standard deviation of the success rate over 5 random seeds. SkiLD outperforms all baselines for most tasks, converging faster and to higher returns.

without skills struggles to utilize that information in downstream tasks. DIAYN learns skills, but few manipulate the objects, so a downstream model struggles to utilize those skills to achieve meaningful rewards. By comparison, SkiLD achieves superior performance on 9 of the 10 downstream tasks evaluated. In the two hardest tasks which require a very long sequence of precise controls, Clean Rag and Cut Peach, SkiLD is the only method that can achieve a non-zero success rate (although still far from fully mastering the tasks), showcasing the potential of local dependencies for skill learning.

#### 4.5 Graph and Diversity Ablations

We also explore the functionality of the graph and diversity components of the skill parameter  $z$  by assessing the downstream performance of SkiLD without these components. This produces two ablative versions of SkiLD: SkiLD without diversity and SkiLD without dependency graphs. To isolate learning from the effect of learned local dependencies, we use ground truth dependency graphs for ablative evaluations where relevant. In Figure 6, learning without graphs results in zero performance, consistent with DIAYN results. In addition, removing diversity produces a notable decline in performance, especially on more challenging tasks like cleaning the rag. These evaluations demonstrate that SkiLD benefits from both the incorporation of dependency graphs and diversity.

## 5 Related Work

This work lies in the unsupervised skill learning framework [35], where the agent must discover a set of useful skills which are reward independent. It then extends these skills to construct a 2-layer hierarchical structure [58], where the upper policy receives reward both for achieving novel skills, and can then be tuned to utilize the learned skills to accomplish an end task. Finally, the skills are identified using token causality, a specific problem identified in causal literature.

### 5.1 Unsupervised Skill Learning

This work describes a framework for utilizing local dependency graphs and diversity to discover unsupervised skills. Diversity-based state coverage skills have been explored in literature [18] utilizing forward and backward mutual information techniques to learn a goal space  $\mathcal{Z}$ , and a skill encoder  $q(z|\cdot)$  [10]. This unsupervised paradigm has been extended with lipschitz constraints [47], contrastive objectives [37], information bottleneck [33], population based methods such as particle estimation [43], quality diversity [42] and mixture of experts [11]. These skills can then be used for hierarchical policies or planners [54, 64, 22], which mirrors the same structure as SkiLD. Unlike these methods, SkiLD adds additional subdivision through dependency graphs, which prevents the combinatorial explosion of skills that can result from trying to cover a large factored space.






Figure 6: A figure illustrating the ablative performance of SkiLD without diversity or without graphs. Without graphs, the method collapses completely, while removing diversity results in a noticeable reduction in downstream performance.

## 5.2 Hierarchical Reinforcement Learning

The hierarchical policy structure in SkiLD where a higher level policy passes a parameter to be interpreted by low-level planners has been formalized in [58], and learned using deep networks utilizing extrinsic reward [2, 59], attention mechanisms [15], initiation criteria [32, 3] and deliberation cost [25]. Hierarchies of goal-based policies [38] has been extended with object-centric representations [63], offline data [46], empowerment [39] and goal counts [49]. In practice, SkiLD uses graph and diversity parameters similar to goal-based methods. However, the space of goals can often be intractable large, and methods to address this use graph laplacians [34] causal chains [12, 13] or general causal relationships [27]. SkiLD is similar to these causal methods but utilizes local dependence along with general two-layer architectures, thus showing increased generalizability.

## 5.3 Causality in Reinforcement Learning

This work investigates the application of local dependency to hierarchical reinforcement learning. This kind of reasoning has been described as “local causality” or “interactions” in prior RL work for data augmentation [51, 52], learning skill chains [12, 13] and exploration [60]. This work is the first synthesis of unsupervised skill learning and local dependencies applied to general 2-layer hierarchical reinforcement learning. Other general causality work investigates action-influence detection [56, 26], affordance learning [9], model learning [28, 20], critical state identification [44], and disentanglement [16]. In the context of relating local dependency and causal inference, we provide a discussion in Appendix C. SkiLD incorporates causality-inspired local dependence to skill learning, resulting in a robust set of transferable skills.

## 6 Conclusion

Unsupervised skill discovery is a powerful tool for learning useful skills in long-horizon sparse reward tasks. However, many unsupervised skill-learning methods do not take advantage of factored environments, resulting in poor performance in complex environments with several objects. Skill Discovery from Local Dependencies utilizes state-specific dependency graphs, identified using learned pointwise conditional mutual information models, to guide skill discovery. The framework of defining skills according to a dependency graph and diversity goal, combined with a learned sampling scheme, achieves difficult downstream tasks. In domains where typically hand-coded primitive skills are given to the agent like Mini-behavior and Interactive Gibson, SkiLD can achieve high performance without requiring explicit domain knowledge. These impressive results arise intuitively from incorporating local dependencies as skill targets, illuminating a meaningful direction for unsupervised skill learning to be applied to a wider array of environments.

**Limitations and Future Work** An important assumption of SkiLD is its access to factored state space. While factored state space can often be naturally obtained from existing RL benchmarks and many real-world environments, developments in disentangled representation learning [45, 29] will help with extending SkiLD to unfactored image domains. Secondly, SkiLD requires accurate detection of local dependencies. While off-the-shelf methods [60, 56] work well for detecting local dependencies in our experiments, future works that can more accurately detect local dependencies will be beneficial to the performance of SkiLD.

## References

- [1] Marcin Andrychowicz, Filip Wolski, Alex Ray, Jonas Schneider, Rachel Fong, Peter Welinder, Bob McGrew, Josh Tobin, OpenAI Pieter Abbeel, and Wojciech Zaremba. Hindsight experience replay. *Advances in neural information processing systems*, 30, 2017.
- [2] Pierre-Luc Bacon, Jean Harb, and Doina Precup. The option-critic architecture. In *Proceedings of the AAAI conference on artificial intelligence*, volume 31, 2017.
- [3] Akhil Bagaria, Ben Abbatematteo, Omer Gottesman, Matt Corsaro, Sreehari Rammohan, and George Konidaris. Effectively learning initiation sets in hierarchical reinforcement learning. *Advances in Neural Information Processing Systems*, 36, 2024.
- [4] Sander Beckers, Hana Chockler, and Joseph Halpern. A causal analysis of harm. *Advances in Neural Information Processing Systems*, 35:2365–2376, 2022.
- [5] Gianluca Bontempi and Maxime Flauder. From dependency to causality: a machine learning approach. *J. Mach. Learn. Res.*, 16(1):2437–2457, 2015.
- [6] Serena Booth, W Bradley Knox, Julie Shah, Scott Niekum, Peter Stone, and Alessandro Allievi. The perils of trial-and-error reward design: misdesign through overfitting and invalid task specifications. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 37, pages 5920–5929, 2023.
- [7] Craig Boutilier, Thomas Dean, and Steve Hanks. Decision-theoretic planning: Structural assumptions and computational leverage. *Journal of Artificial Intelligence Research*, 11:1–94, 1999.
- [8] Craig Boutilier, Nir Friedman, Moises Goldszmidt, and Daphne Koller. Context-specific independence in bayesian networks. *arXiv preprint arXiv:1302.3562*, 2013.
- [9] Jake Brawer, Meiying Qin, and Brian Scassellati. A causal approach to tool affordance learning. In *2020 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 8394–8399. IEEE, 2020.
- [10] Víctor Campos, Alexander Trott, Caiming Xiong, Richard Socher, Xavier Giró-i Nieto, and Jordi Torres. Explore, discover and learn: Unsupervised discovery of state-covering skills. In *International Conference on Machine Learning*, pages 1317–1327. PMLR, 2020.
- [11] Onur Celik, Dongzhuoran Zhou, Ge Li, Philipp Becker, and Gerhard Neumann. Specializing versatile skill libraries using local mixture of experts. In *Conference on Robot Learning*, pages 1423–1433. PMLR, 2022.
- [12] Caleb Chuck, Supawit Chockchowwat, and Scott Niekum. Hypothesis-driven skill discovery for hierarchical deep reinforcement learning. In *2020 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 5572–5579. IEEE, 2020.
- [13] Caleb Chuck, Kevin Black, Aditya Arjun, Yuke Zhu, and Scott Niekum. Granger-causal hierarchical skill discovery. *arXiv preprint arXiv:2306.09509*, 2023.
- [14] Caleb Chuck, Sankaran Vaidyanathan, Stephen Giguere, Amy Zhang, David Jensen, and Scott Niekum. Automated discovery of functional actual causes in complex environments. *arXiv preprint arXiv:2404.10883*, 2024.
- [15] Raviteja Chunduru and Doina Precup. Attention option-critic. *arXiv preprint arXiv:2201.02628*, 2022.
- [16] Oriol Corcoll and Raul Vicente. Disentangling controlled effects for hierarchical reinforcement learning. In Bernhard Schölkopf, Caroline Uhler, and Kun Zhang, editors, *Proceedings of the First Conference on Causal Learning and Reasoning*, volume 177 of *Proceedings of Machine Learning Research*, pages 178–200. PMLR, 11–13 Apr 2022. URL <https://proceedings.mlr.press/v177/corcoll122a.html>.

- [17] Yuqing Du, Eliza Kosoy, Alyssa Dayan, Maria Rufova, Pieter Abbeel, and Alison Gopnik. What can ai learn from human exploration? intrinsically-motivated humans and agents in open-world exploration. In *NeurIPS 2023 workshop: Information-Theoretic Principles in Cognitive Systems*, 2023.
- [18] Benjamin Eysenbach, Abhishek Gupta, Julian Ibarz, and Sergey Levine. Diversity is all you need: Learning skills without a reward function. *arXiv preprint arXiv:1802.06070*, 2018.
- [19] Alhussein Fawzi, Matej Balog, Aja Huang, Thomas Hubert, Bernardino Romera-Paredes, Mohammadamin Barekatin, Alexander Novikov, Francisco J R Ruiz, Julian Schrittwieser, Grzegorz Swirszcz, et al. Discovering faster matrix multiplication algorithms with reinforcement learning. *Nature*, 610(7930):47–53, 2022.
- [20] Fan Feng and Sara Magliacane. Learning dynamic attribute-factored world models for efficient multi-object reinforcement learning. *Advances in Neural Information Processing Systems*, 36, 2024.
- [21] Pierre Fournier, Cédric Colas, Mohamed Chetouani, and Olivier Sigaud. Clic: Curriculum learning and imitation for object control in nonrewarding environments. *IEEE Transactions on Cognitive and Developmental Systems*, 13(2):239–248, 2019.
- [22] Tuomas Haarnoja, Kristian Hartikainen, Pieter Abbeel, and Sergey Levine. Latent space policies for hierarchical reinforcement learning. In *International Conference on Machine Learning*, pages 1851–1860. PMLR, 2018.
- [23] Joseph Y Halpern. *Actual causality*. MIT Press, 2016.
- [24] Joseph Y Halpern and Judea Pearl. Causes and explanations: A structural-model approach. part i: Causes. *The British journal for the philosophy of science*, 2005.
- [25] Jean Harb, Pierre-Luc Bacon, Martin Klissarov, and Doina Precup. When waiting is not an option: Learning options with a deliberation cost. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 32, 2018.
- [26] Jiaheng Hu, Peter Stone, and Roberto Martín-Martín. Causal policy gradient for whole-body mobile manipulation. *arXiv preprint arXiv:2305.04866*, 2023.
- [27] Xing Hu, Rui Zhang, Ke Tang, Jiaming Guo, Qi Yi, Ruizhi Chen, Zidong Du, Ling Li, Qi Guo, Yunji Chen, et al. Causality-driven hierarchical structure discovery for reinforcement learning. *Advances in Neural Information Processing Systems*, 35:20064–20076, 2022.
- [28] Yixuan Huang, Adam Conkey, and Tucker Hermans. Planning for multi-object manipulation with graph neural network relational classifiers. In *2023 IEEE International Conference on Robotics and Automation (ICRA)*, pages 1822–1829. IEEE, 2023.
- [29] Jindong Jiang, Fei Deng, Gautam Singh, and Sungjin Ahn. Object-centric slot diffusion. *arXiv preprint arXiv:2303.10834*, 2023.
- [30] Emily Jin, Jiaheng Hu, Zhuoyi Huang, Ruohan Zhang, Jiajun Wu, Li Fei-Fei, and Roberto Martín-Martín. Mini-behavior: A procedurally generated benchmark for long-horizon decision-making in embodied ai. *arXiv preprint arXiv:2310.01824*, 2023.
- [31] Elia Kaufmann, Leonard Bauersfeld, Antonio Loquercio, Matthias Müller, Vladlen Koltun, and Davide Scaramuzza. Champion-level drone racing using deep reinforcement learning. *Nature*, 620(7976):982–987, 2023.
- [32] Khimya Khetarpal, Martin Klissarov, Maxime Chevalier-Boisvert, Pierre-Luc Bacon, and Doina Precup. Options of interest: Temporal abstraction with interest functions. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 34, pages 4444–4451, 2020.
- [33] Jaekyeom Kim, Seohong Park, and Gunhee Kim. Unsupervised skill discovery with bottleneck option learning. In Marina Meila and Tong Zhang, editors, *Proceedings of the 38th International Conference on Machine Learning*, volume 139 of *Proceedings of Machine Learning Research*, pages 5572–5582. PMLR, 18–24 Jul 2021. URL <https://proceedings.mlr.press/v139/kim21j.html>.

- [34] Martin Klissarov and Marlos C. Machado. Deep Laplacian-based options for temporally-extended exploration. In Andreas Krause, Emma Brunskill, Kyunghyun Cho, Barbara Engelhardt, Sivan Sabato, and Jonathan Scarlett, editors, *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 17198–17217. PMLR, 23–29 Jul 2023. URL <https://proceedings.mlr.press/v202/klissarov23a.html>.
- [35] Pawel Ladosz, Lilian Weng, Minwoo Kim, and Hyondong Oh. Exploration in deep reinforcement learning: A survey. *Information Fusion*, 85:1–22, 2022. ISSN 1566-2535. doi: <https://doi.org/10.1016/j.inffus.2022.03.003>. URL <https://www.sciencedirect.com/science/article/pii/S1566253522000288>.
- [36] Michael Laskin, Denis Yarats, Hao Liu, Kimin Lee, Albert Zhan, Kevin Lu, Catherine Cang, Lerrel Pinto, and Pieter Abbeel. Urlb: Unsupervised reinforcement learning benchmark, 2021.
- [37] Michael Laskin, Hao Liu, Xue Bin Peng, Denis Yarats, Aravind Rajeswaran, and Pieter Abbeel. Cic: Contrastive intrinsic control for unsupervised skill discovery. *arXiv preprint arXiv:2202.00161*, 2022.
- [38] Andrew Levy, Robert Platt, and Kate Saenko. Hierarchical reinforcement learning with hindsight. In *International Conference on Learning Representations*, 2019. URL <https://openreview.net/forum?id=ryzECoAcY7>.
- [39] Andrew Levy, Sreehari Rammohan, Alessandro Allievi, Scott Niekum, and George Konidaris. Hierarchical empowerment: Towards tractable empowerment-based skill-learning. *arXiv preprint arXiv:2307.02728*, 2023.
- [40] Chengshu Li, Fei Xia, Roberto Martín-Martín, Michael Lingelbach, Sanjana Srivastava, Bokui Shen, Kent Vainio, Cem Gokmen, Gokul Dharan, Tanish Jain, Andrey Kurenkov, C. Karen Liu, Hyowon Gweon, Jiajun Wu, Li Fei-Fei, and Silvio Savarese. igibson 2.0: Object-centric simulation for robot learning of everyday household tasks, 2021.
- [41] Chengshu Li, Fei Xia, Roberto Martín-Martín, Michael Lingelbach, Sanjana Srivastava, Bokui Shen, Kent Elliott Vainio, Cem Gokmen, Gokul Dharan, Tanish Jain, Andrey Kurenkov, Karen Liu, Hyowon Gweon, Jiajun Wu, Li Fei-Fei, and Silvio Savarese. igibson 2.0: Object-centric simulation for robot learning of everyday household tasks. In Aleksandra Faust, David Hsu, and Gerhard Neumann, editors, *Proceedings of the 5th Conference on Robot Learning*, volume 164 of *Proceedings of Machine Learning Research*, pages 455–465. PMLR, 08–11 Nov 2022. URL <https://proceedings.mlr.press/v164/li22b.html>.
- [42] Bryan Lim, Luca Grillotti, Lorenzo Bernasconi, and Antoine Cully. Dynamics-aware quality-diversity for efficient learning of skill repertoires. In *2022 International Conference on Robotics and Automation (ICRA)*, pages 5360–5366. IEEE, 2022.
- [43] Hao Liu and Pieter Abbeel. Behavior from the void: Unsupervised active pre-training. *Advances in Neural Information Processing Systems*, 34:18459–18473, 2021.
- [44] Haozhe Liu, Mingchen Zhuge, Bing Li, Yuhui Wang, Francesco Faccio, Bernard Ghanem, and Jürgen Schmidhuber. Learning to identify critical states for reinforcement learning from videos. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 1955–1965, 2023.
- [45] Francesco Locatello, Ben Poole, Gunnar Rätsch, Bernhard Schölkopf, Olivier Bachem, and Michael Tschannen. Weakly-supervised disentanglement without compromises. In *International Conference on Machine Learning*, pages 6348–6359. PMLR, 2020.
- [46] Ofir Nachum, Shixiang Shane Gu, Honglak Lee, and Sergey Levine. Data-efficient hierarchical reinforcement learning. *Advances in neural information processing systems*, 31, 2018.
- [47] Seohong Park, Jongwook Choi, Jaekyeom Kim, Honglak Lee, and Gunhee Kim. Lipschitz-constrained unsupervised skill discovery. In *International Conference on Learning Representations*, 2021.

- [48] Seohong Park, Kimin Lee, Youngwoon Lee, and Pieter Abbeel. Controllability-aware unsupervised skill discovery. *arXiv preprint arXiv:2302.05103*, 2023.
- [49] Shubham Pateria, Budhitama Subagdja, Ah-Hwee Tan, and Chai Quek. End-to-end hierarchical reinforcement learning with integrated subgoal discovery. *IEEE Transactions on Neural Networks and Learning Systems*, 33(12):7778–7790, 2021.
- [50] Judea Pearl. *Causality*. Cambridge university press, 2009.
- [51] Silviu Pitis, Elliot Creager, and Animesh Garg. Counterfactual data augmentation using locally factored dynamics. *Advances in Neural Information Processing Systems*, 33:3976–3990, 2020.
- [52] Silviu Pitis, Elliot Creager, Ajay Mandlekar, and Animesh Garg. Mocoda: Model-based counterfactual data augmentation. *Advances in Neural Information Processing Systems*, 35: 18143–18156, 2022.
- [53] David Poole and Nevin Lianwen Zhang. Exploiting contextual independence in probabilistic inference. *Journal of Artificial Intelligence Research*, 18:263–313, 2003.
- [54] Rafael Rodriguez-Sanchez and George Konidaris. Learning abstract world models for value-preserving planning with options. In *NeurIPS 2023 Workshop on Generalization in Planning*, 2023.
- [55] John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms, 2017.
- [56] Maximilian Seitzer, Bernhard Schölkopf, and Georg Martius. Causal influence detection for improving efficiency in reinforcement learning. *Advances in Neural Information Processing Systems*, 34:22905–22918, 2021.
- [57] Wonil Song, Sangryul Jeon, Hyesong Choi, Kwanghoon Sohn, and Dongbo Min. Learning disentangled skills for hierarchical reinforcement learning through trajectory autoencoder with weak labels. *Expert Systems with Applications*, page 120625, 2023.
- [58] Richard S Sutton, Doina Precup, and Satinder Singh. Between mdps and semi-mdps: A framework for temporal abstraction in reinforcement learning. *Artificial intelligence*, 112(1-2): 181–211, 1999.
- [59] Alexander Sasha Vezhnevets, Simon Osindero, Tom Schaul, Nicolas Heess, Max Jaderberg, David Silver, and Koray Kavukcuoglu. Feudal networks for hierarchical reinforcement learning. In *International Conference on Machine Learning*, pages 3540–3549. PMLR, 2017.
- [60] Zizhao Wang, Jiaheng Hu, Peter Stone, and Roberto Martín-Martín. Elden: Exploration via local dependencies. *Advances in Neural Information Processing Systems*, 36, 2024.
- [61] Jiayi Weng, Huayu Chen, Dong Yan, Kaichao You, Alexis Duburcq, Minghao Zhang, Yi Su, Hang Su, and Jun Zhu. Tianshou: A highly modularized deep reinforcement learning library. *Journal of Machine Learning Research*, 23(267):1–6, 2022. URL <http://jmlr.org/papers/v23/21-1127.html>.
- [62] Peter R Wurman, Samuel Barrett, Kenta Kawamoto, James MacGlashan, Kaushik Subramanian, Thomas J Walsh, Roberto Capobianco, Alisa Devlic, Franziska Eckert, Florian Fuchs, et al. Outracing champion gran turismo drivers with deep reinforcement learning. *Nature*, 602(7896): 223–228, 2022.
- [63] Andrii Zadaianchuk, Maximilian Seitzer, and Georg Martius. Self-supervised visual reinforcement learning with object-centric representations. In *International Conference on Learning Representations*, 2021. URL <https://openreview.net/forum?id=xpplmXCbOw1>.
- [64] Jesse Zhang, Haonan Yu, and Wei Xu. Hierarchical reinforcement learning by discovering intrinsic options. *arXiv preprint arXiv:2101.06521*, 2021.

## A Factored Skills

Learning to reach both a desired graph  $g$  and a diversity parameter  $b$  through primitive actions is challenging. First, different graphs often have substantially different characteristics, with some graphs that are easy to achieve (eg. action→agent), and others that are quite challenging and rare (eg. agent, knife, fruit→fruit). Not only would it be challenging for a single policy to encode all of these behaviors, the diversity parameter notwithstanding, but over-training the frequency at which certain graphs are called might vary significantly. Rather than trying to learn a single monolithic policy, then, we instead structure the skill parameterized policy  $\pi_{\text{skill}}$  as a collection of factored skills:  $\pi_{\text{skill},i}$ , for each factor  $i \in \{1, \dots, N\}$ .

This modification to the policy structure results in three changes: **1)** The upper-level action space passes a single row of the graph  $\mathcal{G}$ , denoted with  $g_i$ , and the desired factor  $i$ . **2)** Instead of achieving an entire graph use the achieved row  $\mathbb{1}[g_{\text{achieved},i} = g_i]$ . **3)** The history of seen graphs  $\mathcal{H}$  is replaced with a history of factored graph rows  $\mathcal{H}_f$ .

Define the history of graph rows as  $\mathcal{H}_f := \{\text{unique}(i, g_{\text{achieved},i} \mid \forall i \in 1, \dots, N \quad \forall g_{\text{achieved}} \in \mathcal{D})\}$ . This takes the unique graph rows from all those seen in previous data. Then the upper policy uses the same historical sampling procedure as with unfactored graphs: the policy samples discretely from the new history, which will by default return  $i, g_i$ , a graph row, and the desired factor. This resolves points **1,3**. Point **2** is addressed by replacing Equation 3 with  $\mathbb{1}[g_{\text{achieved},i} = g_i]$ .

Empirically, we found that without this change, the lower policy rarely learns anything, even simple control of the agent.

## B Environment Details

In this section, we provide a detailed description of the environment, including its semantic stages representing internal progress toward task completion, state space, and action space. We also highlight that while each task consists of multiple semantic stages, agents do not have access to this information.




Figure 7: Environments.

**Installing Printer** As shown in Fig. 7(a), the Installing Printer environment is relatively simple, consisting of 3 factors: the agent, a printer, and a table. The task requires the agent to complete the following **stages**: (1) pick up the printer, (2) bring the printer to and place it on the table, and (3) turn on the printer. The discrete state space consists of (i) the agent’s position and direction, (ii) the positions of the printer and whether it is on or off, and (iii) the position of the table. The discrete action space consists of (i) moving forward, turning left or right, (ii) picking up / placing down the printer, and (iii) turning on / off the printer.

**Thawing** As shown in Fig. 7(b) and Fig. 8(a), the Thawing environment consists of 6 factors: the agent, a sink, a refrigerator, and three frozen objects: fish, olive, and date. Thawing each object requires the agent to complete the following **stages**: (1) move to and open the refrigerator, (2) take the frozen fish out of the refrigerator, (3) put the fish into the sink, and (4) turn on the sink to thaw it. The discrete state space consists of (i) the agent’s position and direction, (ii) the positions of all environment entities, (iii) whether the sink door is turned on, (iv) whether the refrigerator door is opened, and (v) the thawing status of three objects. The discrete action space consists of (i) moving forward, turning left or right, (ii) opening / closing the refrigerator, (iii) turning on / off the sink, and (iv) picking up / placing down each object.

**Cleaning Car** As shown in Fig. 7(c), the Cleaning Car environment consists of 7 factors: the agent, a car, a sink, a bucket, a shelf, a rag, and a piece of soap. Cleaning both the car and the rag requires the agent to complete the following **stages**: (1) take the rag off the shelf, (2) put it in the sink, (3) toggle the sink to soak the rag up, (4) clean the car with the soaked rag, (5) take the soap off the self, and (6) clean the rag with the soap inside the bucket. The discrete state space consists of (i) the agent’s position and direction, (ii) the positions of all environment entities, (iii) whether the sink is turned on, (iv) the soak status of the rag, (v) the cleanness of the rag, and (vi) the cleanness of the car. The discrete action space consists of (i) moving forward, turning left or right, (ii) turning on / off the sink, and (iii) picking up / placing down the rag / soap.

**iGibson** As shown in Fig. 7(d), the iGibson environment consists of 4 factors: the robot, a knife, a peach, and a sink. The robot can do the following things: (1) grasp peach: move close to the peach and grasp it, (3) wash peach: grasp the peach and place it into the sink, (3) grasp knife: move close to the knife and grasp it, (4) cut peach: grasp the knife and use it to cut the peach. The continuous state space consists of (i) the robot’s proprioception, (ii) the poses of all environment entities, and (iii) whether the peach is cut. The continuous action space consists of (i) end-effector position change, (ii) base linear and angular velocity, and (iii) gripper torque (to open/close the gripper).

## C Local Dependencies and Causal Inference

In this work, we define local dependencies according to the state factors  $X = (X^1, \dots, X^N)$  and event of interest  $Y$ , which in the context of an MDP is a subset of the next state factors  $X' = (X'^1, \dots, X'^N)$ . In the factored MDP formulation [7], we assume that  $p$ , the transition dynamics, are represented by a dynamic Bayesian network (DBN) which is a time-directed bipartite graph, with edges only from factors in  $X$  to factors in  $X'$ . In this work, we assume that the underlying ground truth DBN, that is the transition function  $p$ , can be decomposed according to subsets of state factors  $\bar{X}$ , such there exists a  $p^{\bar{X}}(Y = y | \bar{X} = x)$  for every state.

The factored transition dynamics analogizes with causal inference in the following way: If the state factors and next state factors are each assigned a causal variable by adding the assumption that they can be independently intervened on, and each next state variable carries an associated unobserved noise variable  $U^i$ , which we assume is independent of any  $X^k$  not connected to  $X'^j$  and any other next state variable  $X'^j$ , then we can represent the transition dynamics  $p$  with a structural causal model (SCM) [50], a graph connecting the causal variables in  $X$  to the causal variables in  $X'$ .

For a particular outcome variable  $Y$  that is one of the next state causal variables  $X'$ , we can describe local dependence in the RL context according to assumptions about the structural causal model. Represent the non-noise parents of  $Y$  as  $\text{pa}(Y)$ , and the noise parents as  $\text{pa}_U(Y)$ . Under normal causal assumptions, the structural causal model for  $Y$  is a function  $f_Y(\text{pa}(Y), \text{pa}_U(Y)) = Y$ . Define  $\bar{X}$  as a subset of the endogenous parents of  $Y$  and  $\bar{U}$  as an equivalent subset of the noise variables. Further define the values that  $\text{pa}(Y)$ ,  $\text{pa}_U(Y)$ ,  $\bar{X}$ ,  $\bar{U}$  can take on as  $\text{pa}(y)$ ,  $\text{pa}_U(y)$ ,  $\bar{x}$ ,  $\bar{u}$  respectively, and  $(\text{pa}(\mathcal{Y}))$ ,  $\bar{\mathcal{X}}$ ,  $\bar{\mathcal{U}}$  as the set of states the parents of  $Y$ , the variables in  $\bar{X}$  and variables in  $\bar{U}$  can take on respectively.

To formalize local invariance, we add the assumption that  $f_Y$  can be decomposed into a series of functions  $(f_{Y1}(\bar{X}_1 = \bar{x}_1, \bar{U}_1 = \bar{u}_1), \dots, f_{Yk}(\bar{X}_k = \bar{x}, \bar{U}_k = \bar{u}_k))$  and  $g_Y(\text{pa}(Y) = \text{pa}(y), \text{pa}_U(Y) = \text{pa}_U(y))$ , where each  $f_{Yi} : \bar{\mathcal{X}} \times \bar{\mathcal{U}} \rightarrow \mathcal{Y}$  and  $g : \text{pa}(\mathcal{Y}) \rightarrow \{1, \dots, k\}$ , a function mapping the parents of  $Y$  to one of the functions. Then if  $f$  is represented as:

$$f(\text{pa}(x), \text{pa}_U(y)) := \sum_{i=1}^k \mathbb{1}(g_Y(\text{pa}(y), \text{pa}_U(y)) = i) f_{Yi}(\bar{x}_i, \bar{u}_i) \quad (5)$$

The local dependence of  $Y = y$  in a particular state  $(x, x')$  is then the set of variables in  $\bar{X}_i$  for the particular  $i$  where  $\mathbb{1}(g_Y(\text{pa}(y), \text{pa}_U(y)) = i) = 1$ , and the pCMI test is a way of uncovering these local dependencies from observational data.

Local dependence has been investigated in the field of context-specific independence [53, 8], which seeks to find particular assignments of a subset of the causal variables under which an outcome is independent of some subset of the inputs. In particular, context-set specific independence [8] determines if a variable is independent of other variables on a particular subset of states, described as

Table 1: Parameters of Skill Learning and Task Learning. Parameters shared if not specified.

	Name	Environments			
		Printer	Thawing	Cleaning Car	iGibson
Skill Policy	algorithm		Rainbow		TD3
	n step		3		5
	skill horizon		30		100
	exploration noise		0.4		0.2
	optimizer		Adam		
	learning rate		$3 \times 10^{-4}$		
Graph Selection Policy	batch size		64		
	algorithm		PPO		
	optimizer		Adam		
	learning rate		$1 \times 10^{-4}$		
	batch size		1024		
	clip ratio		0.1		
	MLP size		[512, 512]		
	GAE $\lambda$		0.95		
Learned Dynamics Model	entropy coefficient		0.1		
	optimizer		Adam		
	learning rate		$3 \times 10^{-4}$		
	batch size		128		
	number of attention layers		1		
	attention embedding size		128		
Task Skill Selection Policy	number of heads		4		
	algorithm		PPO		
	optimizer		Adam		
	learning rate		$1 \times 10^{-4}$		
	batch size		1024		
	clip ratio		0.1		
	MLP size		[512, 512]		
	GAE $\lambda$		0.95		
Training	entropy coefficient		0.02		
	# of random seeds		5		
	diversity reward coefficient $\beta$		0.5		

the partial context set. While our work uses the pCMI test described in Equation 2, context-specific independence focuses on complete independence using knowledge of the structural model.

Alternatively, interactions can be viewed as the causes ( $\bar{X}$ ) of particular effects ( $Y$ ), which have also been investigated under the description of token or actual cause [24] (as opposed to general cause). Actual cause utilizes a series of counterfactual tests to determine if a cause is necessary, sufficient, and minimal for an outcome. Actual cause has primarily been applied in simple, discrete examples [4, 23], making it difficult to directly apply to RL. However, recent work has incorporated the notion of context-specific independence and extended actual cause to more complex domains [14].

## D Implementation Details

The hyperparameters of skill learning and task learning can be found in Table 1. As it is challenging to identify local dependencies using learned dynamics models in Thawing and iGibson environments, we use ground truth local dependencies from simulator. The codebase is built on tianshou [61] for backend RL, though with significant modifications.


The 5 seeds selected are 0 - 4. The experiments were conducted on machines of the following configurations:

- Nvidia A40 GPU; Intel(R) Xeon(R) Gold 6342 CPU @2.80GHz
- Nvidia A100 GPU; Intel(R) Xeon(R) Gold 6342 CPU @2.80GHz




## 658 **E Skill Visualizations**

659 In Figure 8 we visualize three challenging long-horizon skills learned by SkiLD: thawing the olive,  
660 cleaning the car, and cutting the peach. All of these skills require a sequence of interactions that  
661 is difficult to recover without directed behavior. Thus, comparable baselines do not learn skills of  
662 similar complexity. More skill visualizations can be found at: <https://sites.google.com/view/skild>.



(a) Thaw Olive Skill



(b) Clean Car Skill



(c) Cut Fruit Skill

Figure 8: Policy rollouts for learned policies that achieve long horizon tasks **(a)** Mini-BH thaw olive, **(b)** Mini-BH clean car, **(b)** iGibson cut peach.

## NeurIPS Paper Checklist

### 1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: We make the following claims in the abstract and introduction.

- (a) *Existing methods that reach diverse, distinguishable states can struggle in complex environments with state factors because they focus on simple state factors:* This result is empirically verified in the performance results in experimental domains in Figure 5, where DIAYN and CSD struggle to capture any of the skill requiring the manipulation of multiple state factors, such as cleaning the car or cutting the peach. In Figure 4 we also demonstrate how CSD and DIAYN fail to capture graphs that cover multi-factor interactions.
- (b) *Guiding skill learning with interactions and exploring state coverage within that space can uncover desirable interactions:* This result is empirically verified in the performance results of SkiLD illustrated in Figure 5, where SkiLD is able to handle complex, multi-object interactions such as cleaning the rag after cleaning the car or cutting the peach. Figure 4 also illustrates how desirable interactions like putting soap in the bucket or cleaning with the rag are uncovered.
- (c) *The skills learned directed by local dependencies can be reused to efficiently maximize task rewards:* Again, Figure 5 illustrates that compared to vanilla PPO learning, SkiLD, which uses local dependencies, is able to reuse skills with greater sample efficiency, and often even higher downstream performance. In these results, the pretrained skills are reused across tasks in the same domain, such as thawing the fish, thawing the olive and thawing the date. We provide details about the training procedure in Section 4.4.
- (d) *Combining local dependencies with state diversity is important to induce diverse interactions:* This claim is supported in Section 4.5 and Figure 6, where performance noticeably declines after removing the local dependencies and state diversity individually, especially for more long horizon multi-interaction tasks such as cleaning the car or cleaning the rag.

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

### 2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: We explicitly discuss the limitations of this work in the second paragraph of the Conclusion section.

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.

- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

### 3. Theory Assumptions and Proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [NA]

Justification: We do not make any theoretical claims, since this work primarily focuses on incorporating an inductive bias into unsupervised skill learning and empirical assessment of that claim.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

### 4. Experimental Result Reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Section 3 provides a complete description of the method, with a diagram illustrating the dataflow in Figure 2. The environments are described in Section 4.1 with additional details in Appendix B. The number of seeds and amounts of data used for the skill pretraining and task learning is provided in Section 4.4. We provide all hyperparameters in Appendix D.

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.

- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
  - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
  - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
  - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

## 5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: In <https://sites.google.com/view/skild>, we provide all code for the environments, as well as install dependencies for training SkiLD in the Readme of the interactionHRL folder. With the environments, any data for this environment is generated. We also provide a description of the compute requirements for running these experiments in Appendix D.

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

## 6. Experimental Setting/Details

Question: Does the paper specify all the training and test details (e.g., data splits, hyper-parameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: The core experimental setting is presented in Sec. 4. The detailed settings are included in the appendix.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

## 7. Experiment Statistical Significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: We provided error bars in all quantitative results, which are standard deviations computed from 5 random seeds.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

## 8. Experiments Compute Resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: We specify the computation resources used in our experiments in the appendix.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.

- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

## 9. Code Of Ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: We did not violate the NeurIPS Code of Ethics in this research.

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

## 10. Broader Impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [NA]

Justification: To the best of our knowledge, the work performed does not have notable societal impacts that need to be explicitly addressed.

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

## 11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [NA]

Justification: There is no such risk.

Guidelines:



- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

## 12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [\[Yes\]](#)

Justification: We cited the original paper that produced the code package or dataset.

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, [paperswithcode.com/datasets](https://paperswithcode.com/datasets) has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

## 13. New Assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [\[NA\]](#)

Justification: This paper does not release new assets.

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

## 14. Crowdsourcing and Research with Human Subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [\[NA\]](#)



978 Justification: The paper does not involve crowdsourcing nor research with human subjects.

979 Guidelines:

- 980 • The answer NA means that the paper does not involve crowdsourcing nor research with
- 981 human subjects.
- 982 • Including this information in the supplemental material is fine, but if the main contribu-
- 983 tion of the paper involves human subjects, then as much detail as possible should be
- 984 included in the main paper.
- 985 • According to the NeurIPS Code of Ethics, workers involved in data collection, curation,
- 986 or other labor should be paid at least the minimum wage in the country of the data
- 987 collector.

988 **15. Institutional Review Board (IRB) Approvals or Equivalent for Research with Human**

989 **Subjects**

990 Question: Does the paper describe potential risks incurred by study participants, whether

991 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)

992 approvals (or an equivalent approval/review based on the requirements of your country or

993 institution) were obtained?

994 Answer: [NA]

995 Justification: The paper does not involve crowdsourcing nor research with human subjects.

996 Guidelines:

- 997 • The answer NA means that the paper does not involve crowdsourcing nor research with
- 998 human subjects.
- 999 • Depending on the country in which research is conducted, IRB approval (or equivalent)
- 1000 may be required for any human subjects research. If you obtained IRB approval, you
- 1001 should clearly state this in the paper.
- 1002 • We recognize that the procedures for this may vary significantly between institutions
- 1003 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
- 1004 guidelines for their institution.
- 1005 • For initial submissions, do not include any information that would break anonymity (if
- 1006 applicable), such as the institution conducting the review.